

# Laporan Praktikum 1 AMP

Antonius Aditya Rizky Wijaya

G5402221003

2025-01-23

```
setwd("//sapi/HOMEMHS/AKT2022/g5402221003/Analisis Model Prediktif")
```

## Basic Commands

```
x <- c(1, 3, 2, 5)
x
## [1] 1 3 2 5
```

Note that the `>` is not part of the command; rather, it is printed by R to indicate that it is ready for another command to be entered. We can also save things using `=` rather than `<-`:

```
x = c(1, 6, 2)
x
## [1] 1 6 2
y = c(1, 4, 3)
```

Hitting the *up* arrow multiple times will display the previous commands, which can then be edited. This is useful since one often wishes to repeat a similar command. In addition, typing `?funcname` will always cause R to open a new help file window with additional information about the function `funcname()`.

We can tell R to add two sets of numbers together. It will then add the first number from `x` to the first number from `y`, and so on. However, `x` and `y` should be the same length. We can check their length using the `length()` function.

```
length(x)
## [1] 3
length(y)
## [1] 3
x + y
## [1] 2 10 5
```

The `ls()` function allows us to look at a list of all of the objects, such as data and functions, that we have saved so far. The `rm()` function can be used to delete any that we don't want.

```
ls()
## [1] "x" "y"
rm(x, y)
ls()
## character(0)
```

It's also possible to remove all objects at once:

```
rm(list = ls())
```

The `matrix()` function can be used to create a matrix of numbers. Before we use the `matrix()` function, we can learn more about it:

```
?matrix
## starting httpd help server ... done
```

The help file reveals that the `matrix()` function takes a number of inputs, but for now we focus on the first three: the data (the entries in the matrix), the number of rows, and the number of columns. First, we create a simple matrix.

```
x <- matrix(data = c(1, 2, 3, 4), nrow = 2, ncol = 2)
x
##      [,1] [,2]
## [1,]    1    3
## [2,]    2    4
```

Note that we could just as well omit typing `data=`, `nrow=`, and `ncol=` in the `matrix()` command above: that is, we could just type

```
x <- matrix(c(1, 2, 3, 4), 2, 2)
```

and this would have the same effect. However, it can sometimes be useful to specify the names of the arguments passed in, since otherwise R will assume that the function arguments are passed into the function in the same order that is given in the function's help file. As this example illustrates, by default R creates matrices by successively filling in columns. Alternatively, the `byrow = TRUE` option can be used to populate the matrix in order of the rows.

```
matrix(c(1, 2, 3, 4), 2, 2, byrow = TRUE)
##      [,1] [,2]
## [1,]    1    2
## [2,]    3    4
```

Notice that in the above command we did not assign the matrix to a value such as `x`. In this case the matrix is printed to the screen but is not saved for future calculations. The `sqrt()` function returns the square root of each element of a vector or matrix. The command `x^2`

raises each element of `x` to the power 2; any powers are possible, including fractional or negative powers.

```
sqrt(x)

##           [,1]      [,2]
## [1,]  1.000000  1.732051
## [2,]  1.414214  2.000000

x^2

##           [,1] [,2]
## [1,]         1    9
## [2,]         4   16
```

The `rnorm()` function generates a vector of random normal variables, with first argument `n` the sample size. Each time we call this function, we will get a different answer. Here we create two correlated sets of numbers, `x` and `y`, and use the `cor()` function to compute the correlation between them.

```
x <- rnorm(50)
y <- x + rnorm(50, mean = 50, sd = .1)
cor(x, y)

## [1] 0.996153
```

By default, `rnorm()` creates standard normal random variables with a mean of 0 and a standard deviation of 1. However, the mean and standard deviation can be altered using the `mean` and `sd` arguments, as illustrated above. Sometimes we want our code to reproduce the exact same set of random numbers; we can use the `set.seed()` function to do this. The `set.seed()` function takes an (arbitrary) integer argument.

```
set.seed(1303)
rnorm(50)

## [1] -1.1439763145  1.3421293656  2.1853904757  0.5363925179  0.0631929665
## [6]  0.5022344825 -0.0004167247  0.5658198405 -0.5725226890 -1.1102250073
## [11] -0.0486871234 -0.6956562176  0.8289174803  0.2066528551 -0.2356745091
## [16] -0.5563104914 -0.3647543571  0.8623550343 -0.6307715354  0.3136021252
## [21] -0.9314953177  0.8238676185  0.5233707021  0.7069214120  0.4202043256
## [26] -0.2690521547 -1.5103172999 -0.6902124766 -0.1434719524 -1.0135274099
## [31]  1.5732737361  0.0127465055  0.8726470499  0.4220661905 -0.0188157917
## [36]  2.6157489689 -0.6931401748 -0.2663217810 -0.7206364412  1.3677342065
## [41]  0.2640073322  0.6321868074 -1.3306509858  0.0268888182  1.0406363208
## [46]  1.3120237985 -0.0300020767 -0.2500257125  0.0234144857  1.6598706557
```

We use `set.seed()` throughout the labs whenever we perform calculations involving random quantities. In general this should allow the user to reproduce our results. However, as new versions of R become available, small discrepancies may arise between this book and the output from R.

The `mean()` and `var()` functions can be used to compute the mean and variance of a vector of numbers. Applying `sqrt()` to the output of `var()` will give the standard deviation. Or we can simply use the `sd()` function.

```
set.seed(3)
y <- rnorm(100)
mean(y)

## [1] 0.01103557

var(y)

## [1] 0.7328675

sqrt(var(y))

## [1] 0.8560768


sd(y)

## [1] 0.8560768
```


## Graphics

The `plot()` function is the primary way to plot data in R. For instance, `plot(x, y)` produces a scatterplot of the numbers in `x` versus the numbers in `y`. There are many additional options that can be passed in to the `plot()` function.

```
x <- rnorm(100)
y <- rnorm(100)
plot(x, y)
```



```
plot(x, y, xlab = "this is the x-axis",  
     ylab = "this is the y-axis",  
     main = "Plot of X vs Y")
```



We will often want to save the output of an R plot. The command that we use to do this will depend on the file type that we would like to create. For instance, to create a pdf, we use the `pdf()` function, and to create a jpeg, we use the `jpeg()` function.

```
pdf("Figure.pdf")
plot(x, y, col = "green")
dev.off()
```

```
## png
## 2
```

The function `dev.off()` indicates to R that we are done creating the plot. Alternatively, we can simply copy the plot window and paste it into an appropriate file type, such as a Word document.

The function `seq()` can be used to create a sequence of numbers. For instance, `seq(a, b)` makes a vector of integers between `a` and `b`. There are many other options: for instance, `seq(0, 1, length = 10)` makes a sequence of 10 numbers that are equally spaced between 0 and 1. Typing `3:11` is a shorthand for `seq(3, 11)` for integer arguments.

```
x <- seq(1, 10)
x
## [1] 1 2 3 4 5 6 7 8 9 10

x <- 1:10
x
## [1] 1 2 3 4 5 6 7 8 9 10



x <- seq(-pi, pi, length = 50)
```

We will now create some more sophisticated plots. The `contour()` function produces a *contour plot* in order to represent three-dimensional data; it is like a topographical map. It takes three arguments:

- A vector of the `x` values (the first dimension),
- A vector of the `y` values (the second dimension), and
- A matrix whose elements correspond to the `z` value (the third dimension) for each pair of `(x, y)` coordinates.


As with the `plot()` function, there are many other inputs that can be used to fine-tune the output of the `contour()` function. To learn more about these, take a look at the help file by typing `?contour`.

```
y <- x
f <- outer(x, y, function(x, y) cos(y) / (1 + x^2))
contour(x, y, f)
contour(x, y, f, nlevels = 45, add = T)
```




The `image()` function works the same way as `contour()`, except that it produces a color-coded plot whose colors depend on the  $z$  value. This is known as a *heatmap*, and is sometimes used to plot temperature in weather forecasts. Alternatively, `persp()` can be used to produce a three-dimensional plot. The arguments `theta` and `phi` control the angles at which the plot is viewed.

```
image(x, y, fa)
```




```
persp(x, y, fa)
```







```
persp(x, y, fa, theta = 30)
```




```
persp(x, y, fa, theta = 30, phi = 20)
```



```
persp(x, y, fa, theta = 30, phi = 70)
```



```
persp(x, y, fa, theta = 30, phi = 40)
```



## Indexing Data

We often wish to examine part of a set of data. Suppose that our data is stored in the matrix A.

```
A <- matrix(1:16, 4, 4)
A
##      [,1] [,2] [,3] [,4]
## [1,]    1    5    9   13
## [2,]    2    6   10   14
## [3,]    3    7   11   15
## [4,]    4    8   12   16
```

Then, typing

```
A[2, 3]
## [1] 10
```

will select the element corresponding to the second row and the third column. The first number after the open-bracket symbol [ always refers to the row, and the second number always refers to the column. We can also select multiple rows and columns at a time, by providing vectors as the indices.

```
A[c(1, 3), c(2, 4)]
```

```
##      [,1] [,2]
## [1,]    5  13
## [2,]    7  15
```

```
A[1:3, 2:4]
```

```
##      [,1] [,2] [,3]
## [1,]    5    9  13
## [2,]    6   10  14
## [3,]    7   11  15
```

```
A[1:2, ]
```

```
##      [,1] [,2] [,3] [,4]
## [1,]    1    5    9  13
## [2,]    2    6   10  14
```

```
A[, 1:2]
```

```
##      [,1] [,2]
## [1,]    1    5
## [2,]    2    6
## [3,]    3    7
## [4,]    4    8
```

The last two examples include either no index for the columns or no index for the rows. These indicate that R should include all columns or all rows, respectively. R treats a single row or column of a matrix as a vector.

```
A[1, ]
```

```
## [1]  1  5  9 13
```

The use of a negative sign - in the index tells R to keep all rows or columns except those indicated in the index.

```
A[-c(1, 3), ]
```

```
##      [,1] [,2] [,3] [,4]
## [1,]    2    6   10  14
## [2,]    4    8   12  16
```

```
A[-c(1, 3), -c(1, 3, 4)]
```

```
## [1] 6 8
```

The `dim()` function outputs the number of rows followed by the number of columns of a given matrix.

```
dim(A)
```

```
## [1] 4 4
```

## Loading Data

```
#Auto <- read.table("Auto.data")
#View(Auto)
#head(Auto)
```

Note that Auto.data is simply a text file, which you could alternatively open on your computer using a standard text editor. It is often a good idea to view a data set using a text editor or other software such as Excel before loading it into R.

This particular data set has not been loaded correctly, because R has assumed that the variable names are part of the data and so has included them in the first row. The data set also includes a number of missing observations, indicated by a question mark ?. Missing values are a common occurrence in real data sets. Using the option `header = T` (or `header = TRUE`) in the `read.table()` function tells R that the first line of the file contains the variable names, and using the option `na.strings` tells R that any time it sees a particular character or set of characters (such as a question mark), it should be treated as a missing element of the data matrix.

```
#Auto <- read.table("Auto.data", header = T, na.strings = "?",
stringsAsFactors = T)
#View(Auto)
```

The `stringsAsFactors = T` argument tells R that any variable containing character strings should be interpreted as a qualitative variable, and that each distinct character string represents a distinct level for that qualitative variable. An easy way to load data from Excel into R is to save it as a csv (comma-separated values) file, and then use the `read.csv()` function.

```
Auto <- read.csv("Auto.csv", na.strings = "?", stringsAsFactors = T)
View(Auto)
dim(Auto)

## [1] 397  9

Auto[1:4, ]

##   mpg cylinders displacement horsepower weight acceleration year origin
## 1   18         8          307         130   3504          12.0    70      1
## 2   15         8          350         165   3693          11.5    70      1
## 3   18         8          318         150   3436          11.0    70      1
## 4   16         8          304         150   3433          12.0    70      1
##                                     name
## 1 chevrolet chevelle malibu
## 2          buick skylark 320
## 3          plymouth satellite
## 4             amc rebel sst
```

The `dim()` function tells us that the data has 397 observations, or rows, and nine variables, or columns. There are various ways to deal with the missing data. In this case, only five of

the rows contain missing observations, and so we choose to use the `na.omit()` function to simply remove these rows.

```
Auto <- na.omit(Auto)
dim(Auto)

## [1] 392  9
```

Once the data are loaded correctly, we can use `names()` to check the variable names.

```
names(Auto)

## [1] "mpg"          "cylinders"     "displacement" "horsepower"   "weight"
## [6] "acceleration" "year"         "origin"       "name"
```


## Additional Graphical and Numerical Summaries

We can use the `plot()` function to produce *scatterplots* of the quantitative variables. However, simply typing the variable names will produce an error message, because R does not know to look in the `Auto` data set for those variables.


```
#plot(cylinders, mpg)
```

To refer to a variable, we must type the data set and the variable name joined with a `$` symbol. Alternatively, we can use the `attach()` function in order to tell R to make the variables in this data frame available by name.

```
plot(Auto$cylinders, Auto$mpg)
```



```
attach(Auto)
plot(cylinders, mpg)
```




The cylinders variable is stored as a numeric vector, so R has treated it as quantitative. However, since there are only a small number of possible values for cylinders, one may prefer to treat it as a qualitative variable. The `as.factor()` function converts quantitative variables into qualitative variables.


```
cylinders <- as.factor(cylinders)
```

If the variable plotted on the  $x$ -axis is qualitative, then *boxplots* will automatically be produced by the `plot()` function. As usual, a number of options can be specified in order to customize the plots.

```
plot(cylinders, mpg)
```




```
plot(cylinders, mpg, col = "red")
```




```
plot(cylinders, mpg, col = "red", varwidth = T)
```






```
plot(cylinders, mpg, col = "red", varwidth = T,
     horizontal = T)
```




```
plot(cylinders, mpg, col = "red", varwidth = T,  
     xlab = "cylinders", ylab = "MPG")
```




The `hist()` function can be used to plot a *histogram*. Note that `col = 2` has the same effect as `col = "red"`.


```
hist(mpg)
```



```
hist(mpg, col = 2)
```

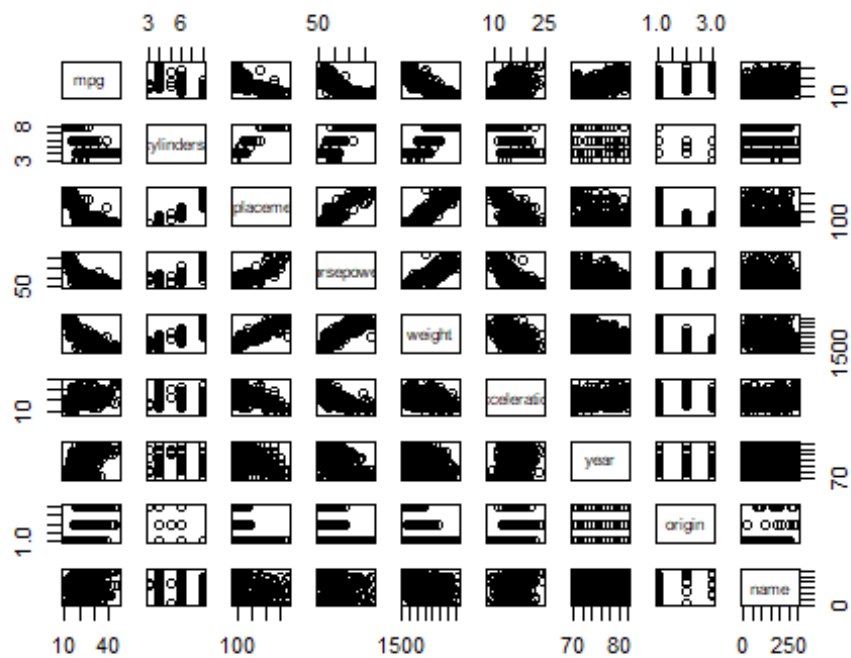


```
hist(mpg, col = 2, breaks = 15)
```



The `pairs()` function creates a *scatterplot matrix*, i.e. a scatterplot for every pair of variables. We can also produce scatterplots for just a subset of the variables.

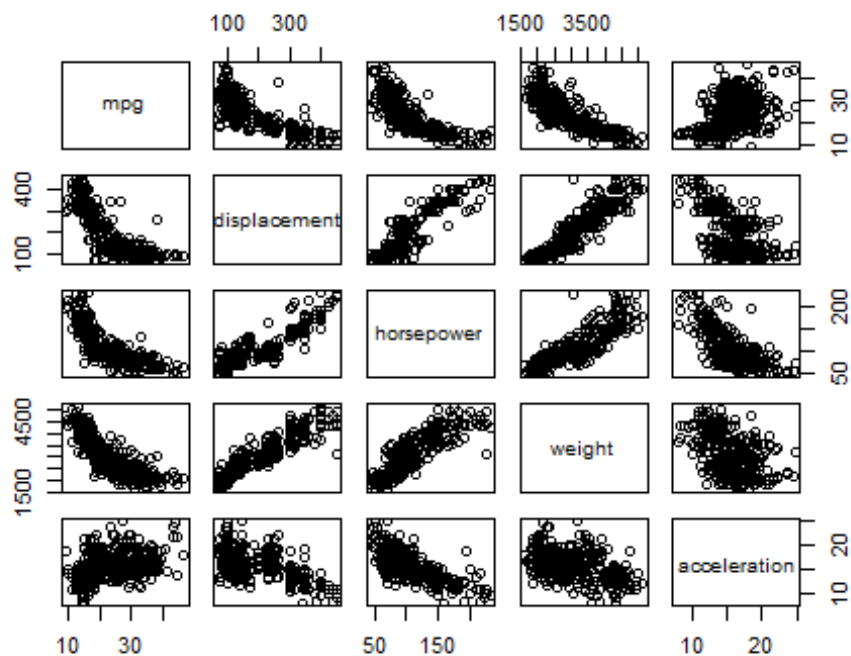
```
pairs(Auto)
```



```


pairs(
  ~ mpg + displacement + horsepower + weight + acceleration,
  data = Auto
)

```



In conjunction with the `plot()` function, `identify()` provides a useful interactive method for identifying the value of a particular variable for points on a plot. We pass in three arguments to `identify()`: the  $x$ -axis variable, the  $y$ -axis variable, and the variable whose values we would like to see printed for each point. Then clicking one or more points in the plot and hitting Escape will cause R to print the values of the variable of interest. The numbers printed under the `identify()` function correspond to the rows for the selected points.

```
plot(horsepower, mpg)
identify(horsepower, mpg, name)
```



```
## integer(0)
```

The `summary()` function produces a numerical summary of each variable in a particular data set.

```
summary(Auto)
```

```
##      mpg      cylinders  displacement  horsepower
weight
## Min.   : 9.00   Min.   :3.000   Min.   : 68.0   Min.   : 46.0   Min.
:1613
## 1st Qu.:17.00   1st Qu.:4.000   1st Qu.:105.0   1st Qu.: 75.0   1st
Qu.:2225
## Median :22.75   Median :4.000   Median :151.0   Median : 93.5   Median
:2804
## Mean   :23.45   Mean   :5.472   Mean   :194.4   Mean   :104.5   Mean
:2978
## 3rd Qu.:29.00   3rd Qu.:8.000   3rd Qu.:275.8   3rd Qu.:126.0   3rd
Qu.:3615
## Max.   :46.60   Max.   :8.000   Max.   :455.0   Max.   :230.0   Max.
:5140
##
##      acceleration      year      origin      name
## Min.   : 8.00   Min.   :70.00   Min.   :1.000   amc matador      : 5
## 1st Qu.:13.78   1st Qu.:73.00   1st Qu.:1.000   ford pinto       : 5
## Median :15.50   Median :76.00   Median :1.000   toyota corolla   : 5
## Mean   :15.54   Mean   :75.98   Mean   :1.577   amc gremlin      : 4
```

```
## 3rd Qu.:17.02 3rd Qu.:79.00 3rd Qu.:2.000 amc hornet : 4
## Max. :24.80 Max. :82.00 Max. :3.000 chevrolet chevette: 4
## (Other) :365
```

For qualitative variables such as name, R will list the number of observations that fall in each category. We can also produce a summary of just a single variable.

```
summary(mpg)
```

```
## Min. 1st Qu. Median Mean 3rd Qu. Max.
## 9.00 17.00 22.75 23.45 29.00 46.60
```

## NOMOR 8

- a. Use the `read.csv()` function to read the data into R. Call the loaded data `college`. Make sure that you have the directory set to the correct location for the data.

```
college <- read.csv("College.csv")
```

- b. Look at the data using the `View()` function. You should notice that the first column is just the name of each university. We don't really want R to treat this as data. However, it may be handy to have these names for later.

```
rownames(college) <- college[, 1]
```

```
View(college)
```

```
college <- college[, -1]
```

```
View(college)
```

- c.
  - i. Use the `summary()` function to produce a numerical summary of the variables in the data set.

```
summary(college)
```

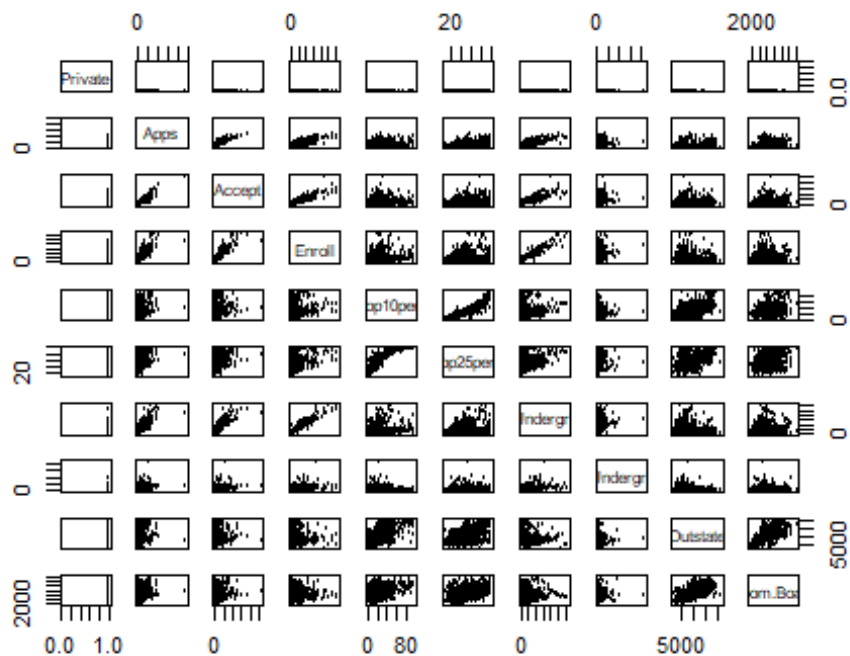
```
## Private Apps Accept Enroll
## Length:777 Min. : 81 Min. : 72 Min. : 35
## Class :character 1st Qu.: 776 1st Qu.: 604 1st Qu.: 242
## Mode :character Median : 1558 Median : 1110 Median : 434
## Mean : 3002 Mean : 2019 Mean : 780
## 3rd Qu.: 3624 3rd Qu.: 2424 3rd Qu.: 902
## Max. :48094 Max. :26330 Max. :6392
## Top10perc Top25perc F.Undergrad P.Undergrad
## Min. : 1.00 Min. : 9.0 Min. : 139 Min. : 1.0
## 1st Qu.:15.00 1st Qu.: 41.0 1st Qu.: 992 1st Qu.: 95.0
## Median :23.00 Median : 54.0 Median : 1707 Median : 353.0
## Mean :27.56 Mean : 55.8 Mean : 3700 Mean : 855.3
## 3rd Qu.:35.00 3rd Qu.: 69.0 3rd Qu.: 4005 3rd Qu.: 967.0
## Max. :96.00 Max. :100.0 Max. :31643 Max. :21836.0
## Outstate Room.Board Books Personal
## Min. : 2340 Min. :1780 Min. : 96.0 Min. : 250
## 1st Qu.: 7320 1st Qu.:3597 1st Qu.: 470.0 1st Qu.: 850
```



```
## Median : 9990    Median :4200    Median : 500.0    Median :1200
## Mean    :10441    Mean    :4358    Mean    : 549.4    Mean    :1341
## 3rd Qu.:12925    3rd Qu.:5050    3rd Qu.: 600.0    3rd Qu.:1700
## Max.    :21700    Max.    :8124    Max.    :2340.0    Max.    :6800
##      PhD      Terminal      S.F.Ratio      perc.alumni
## Min.   : 8.00   Min.   : 24.0   Min.   : 2.50   Min.   : 0.00
## 1st Qu.:62.00   1st Qu.:71.0   1st Qu.:11.50   1st Qu.:13.00
## Median :75.00   Median :82.0   Median :13.60   Median :21.00
## Mean   :72.66   Mean   :79.7   Mean   :14.09   Mean   :22.74
## 3rd Qu.:85.00   3rd Qu.:92.0   3rd Qu.:16.50   3rd Qu.:31.00
## Max.   :103.00   Max.   :100.0   Max.   :39.80   Max.   :64.00
##      Expend      Grad.Rate
## Min.   : 3186   Min.   : 10.00
## 1st Qu.:6751   1st Qu.:53.00
## Median :8377   Median :65.00
## Mean   :9660   Mean   :65.46
## 3rd Qu.:10830  3rd Qu.:78.00
## Max.   :56233  Max.   :118.00
```

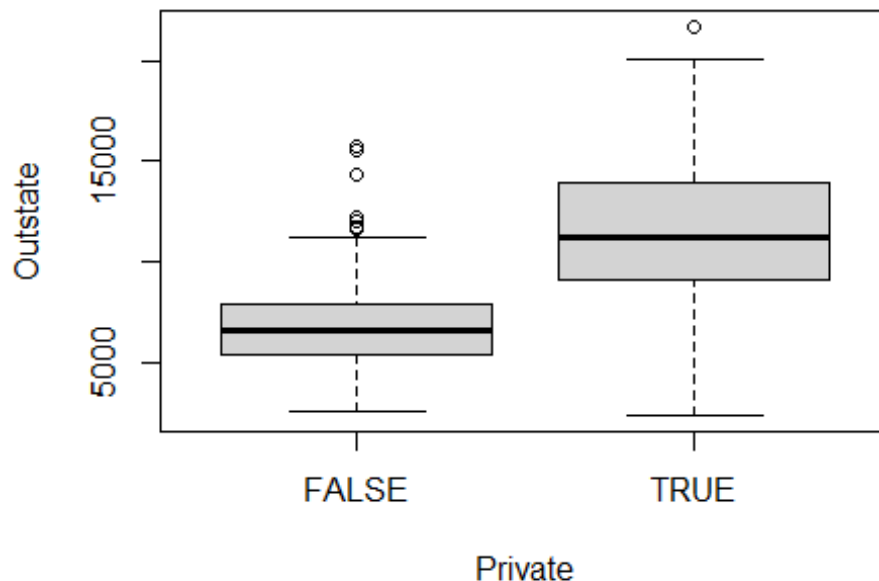
- ii. Use the `pairs()` function to produce a scatterplot matrix of the first ten columns or variables of the data. Recall that you can reference the first ten columns of a matrix A using `A[,1:10]`.

```
college$Private <- college$Private == "Yes"
pairs(college[, 1:10], cex = 0.2)
```



iii. Use the `plot()` function to produce side-by-side boxplots of `Outstate` versus `Private`.

```
plot(college$Outstate ~ factor(college$Private), xlab = "Private", ylab = "Outstate")
```



iv. Create a new qualitative variable, called `Elite`, by binning the `Top10perc` variable. We are going to divide universities into two groups based on whether or not the proportion of students coming from the top 10% of their high school classes exceeds 50%.


```
```\n> Elite <- rep("No", nrow(college))\n> Elite[college$Top10perc > 50] <- "Yes"\n> Elite <- as.factor(Elite)\n> college <- data.frame(college, Elite)\n```\n
```

Use the `summary()` function to see how many elite universities there are. Now use the `plot()` function to produce side-by-side boxplots of `Outstate` versus `Elite`.

```
college$Elite <- factor(ifelse(college$Top10perc > 50, "Yes", "No"))\nsummary(college$Elite)
```

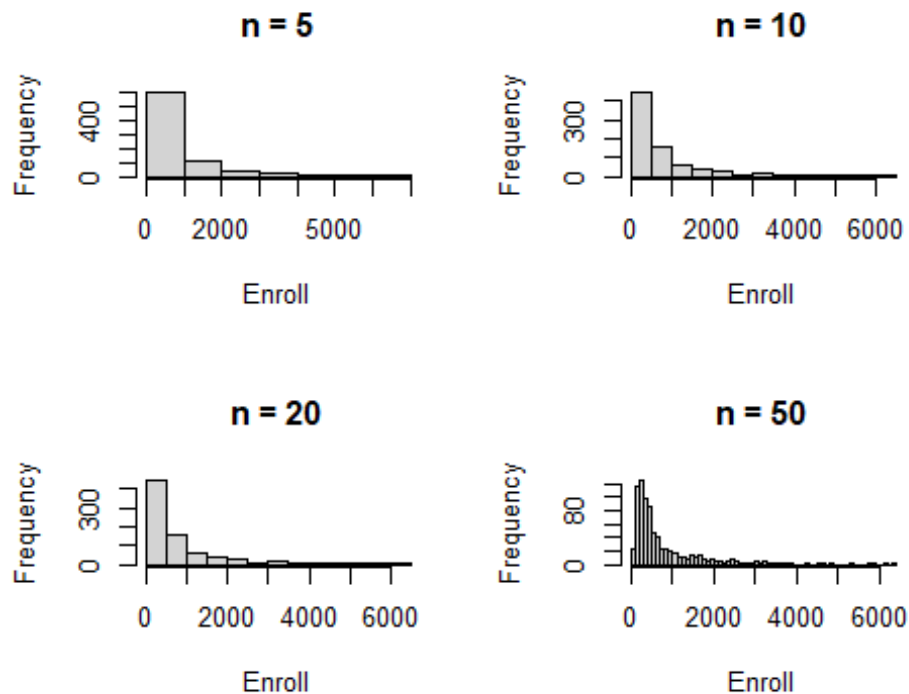
```
## No Yes\n## 699  78
```

```
plot(college$Outstate ~ college$Elite, xlab = "Elite", ylab = "Outstate")
```



- v. Use the `hist()` function to produce some histograms with differing numbers of bins for a few of the quantitative variables. You may find the command `par(mfrow=c(2,2))` useful: it will divide the print window into four regions so that four plots can be made simultaneously. Modifying the arguments to this function will divide the screen in other ways.

```
par(mfrow = c(2, 2))
for (n in c(5, 10, 20, 50)) {
  hist(college$Enroll, breaks = n, main = paste("n =", n), xlab = "Enroll")
}
```



vi. Continue exploring the data, and provide a brief summary of what you discover.

```
chisq.test(college$Private, college$Elite)
```

```
##
## Pearson's Chi-squared test with Yates' continuity correction
##
## data: college$Private and college$Elite
## X-squared = 4.3498, df = 1, p-value = 0.03701
```

## NOMOR 9

a. Which of the predictors are quantitative, and which are qualitative?

```
auto <- read.csv("Auto.csv")
auto <- na.omit(auto)
```

Kuantitatif : mpg, silinder, perpindahan, tenaga kuda, berat, tahun dan percepatan  
Kualitatif : asal dan nama

b. What is the range of each quantitative predictor? You can answer this using the range() function.

```
#apply(auto, 2, range)
sapply(auto[, 1:7], range)
```

```
##      mpg      cylinders displacement horsepower weight acceleration year
## [1,] "9"       "3"         "68"         "?"         "1613" "8"         "70"
## [2,] "46.6"    "8"         "455"         "98"         "5140" "24.8"       "82"
```

c. What is the mean and standard deviation of each quantitative predictor?

```
sapply(auto[, 1:7], mean)

## Warning in mean.default(X[[i]], ...): argument is not numeric or logical:
## returning NA

##      mpg      cylinders displacement      horsepower      weight
acceleration
## 23.515869      5.458438  193.532746              NA 2970.261965
15.555668
##      year
## 75.994962

sapply(auto[, 1:7], sd)

## Warning in var(if (is.vector(x) || is.factor(x)) x else as.double(x),
na.rm =
## na.rm): NAs introduced by coercion

##      mpg      cylinders displacement      horsepower      weight
acceleration
## 7.825804      1.701577  104.379583              NA 847.904119
2.749995
##      year
## 3.690005
```

d. Now remove the 10th through 85th observations. What is the range, mean, and standard deviation of each predictor in the subset of the data that remains?

```
auto_new=auto[-(10:85),]
#Auto[-(10:85),]

apply(auto_new, 2, range)

##      mpg      cylinders displacement horsepower weight acceleration year
origin
## [1,] "11.0" "3"         " 68"         "?"         "1649" " 8.5"       "70" "1"
## [2,] "46.6" "8"         "455"         "98"         "4997" "24.8"       "82" "3"
##      name
## [1,] "amc ambassador brougham"
## [2,] "vw rabbit custom"

sapply(auto_new[, 1:7], mean)

## Warning in mean.default(X[[i]], ...): argument is not numeric or logical:
## returning NA

##      mpg      cylinders displacement      horsepower      weight
acceleration
```

```
##      24.438629      5.370717    187.049844      NA    2933.962617
15.723053
##      year
##      77.152648

sapply(auto_new[, 1:7], sd)

## Warning in var(if (is.vector(x) || is.factor(x)) x else as.double(x),
na.rm =
## na.rm): NAs introduced by coercion

##      mpg      cylinders displacement    horsepower      weight
acceleration
##      7.908184      1.653486     99.635385          NA    810.642938
2.680514
##      year
##      3.111230
```

- e. Using the full data set, investigate the predictors graphically, using scatterplots or other tools of your choice. Create some plots highlighting the relationships among the predictors. Comment on your findings.

```
library(ggplot2)


## Warning: package 'ggplot2' was built under R version 4.3.3

##
## Attaching package: 'ggplot2'

## The following object is masked from 'Auto':
##
##      mpg

ggplot(auto, aes(weight, displacement)) + geom_point() + geom_smooth()

## `geom_smooth()` using method = 'loess' and formula = 'y ~ x'
```



Perpindahan dan berat berkorelasi positif dengan beberapa hubungan bersifat non-linier.

- f. Suppose that we wish to predict gas mileage (mpg) on the basis of the other variables. Do your plots suggest that any of the other variables might be useful in predicting mpg? Justify your answer.

Dari plot di soal e, hampir semua prediktor berkorelasi dengan mpg kecuali nama.

## NOMOR 10

- a. To begin, load in the Boston data set. The Boston data set is part of the ISLR2 library in R.

```
library(ISLR2)
```

```
## Warning: package 'ISLR2' was built under R version 4.3.3
```

```
##
```

```
## Attaching package: 'ISLR2'
```

```
## The following object is masked _by_ '.GlobalEnv':
```

```
##
```

```
##      Auto
```

```
head(Boston)
```

```
##      crim zn indus chas   nox   rm   age   dis rad tax ptratio lstat medv
## 1 0.00632 18  2.31    0 0.538 6.575 65.2 4.0900  1 296   15.3  4.98 24.0
```

```
## 2 0.02731 0 7.07 0 0.469 6.421 78.9 4.9671 2 242 17.8 9.14 21.6
## 3 0.02729 0 7.07 0 0.469 7.185 61.1 4.9671 2 242 17.8 4.03 34.7
## 4 0.03237 0 2.18 0 0.458 6.998 45.8 6.0622 3 222 18.7 2.94 33.4
## 5 0.06905 0 2.18 0 0.458 7.147 54.2 6.0622 3 222 18.7 5.33 36.2
## 6 0.02985 0 2.18 0 0.458 6.430 58.7 6.0622 3 222 18.7 5.21 28.7
```

```
dim(Boston)
```

```
## [1] 506 13
```

```
?Boston
```


How many rows are in this data set? How many columns? What do the rows and columns represent?

Ada 506 baris dan 14 kolom. Setiap baris mewakili satu rumah beserta atributnya. Setiap kolom mewakili satu set atribut rumah.

- b. Make some pairwise scatterplots of the predictors (columns) in this data set. Describe your findings.

```
library(ggplot2)
```

```
ggplot(Boston, aes(rad, crim)) + geom_point()
```



Semakin tinggi indeks akses ke jalan raya radial, semakin tinggi pula kejahatan.

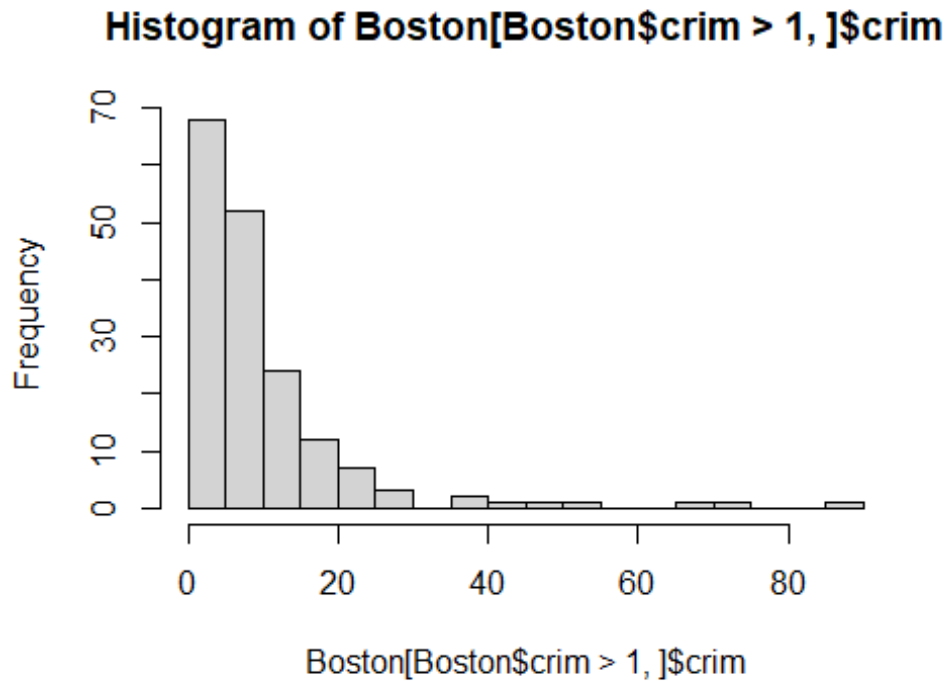
- c. Are any of the predictors associated with per capita crime rate? If so, explain the relationship.




Ya, semakin tinggi indeks akses ke jalan raya radial, semakin tinggi pula kejahatan.

- d. Do any of the census tracts of Boston appear to have particularly high crime rates? Tax rates? Pupil-teacher ratios? Comment on the range of each predictor.


```
hist(Boston[Boston$crim>1,]$crim, breaks=25)
```



```
hist(Boston$tax, breaks=25)
```



```
hist(Boston$ptratio, breaks=25)
```



Tingkat kejahatan :

sebagian besar kota memiliki tingkat kejahatan yang rendah, tetapi beberapa daerah pinggiran kota memiliki tingkat kejahatan lebih dari 20 bahkan ada yang diatas 80.

Pajak : ada kesenjangan antara kota berpajak rendah dengan yang tinggi, puncaknya sekitar 0,5%.

Rasio murid-guru : tidak seimbang

e. How many of the census tracts in this data set bound the Charles river?

```
sum(Boston$chas)
```

```
## [1] 35
```

Ada 35 daerah pinggiran kota yang berbatasan dengan sungai Charles.

f. What is the median pupil-teacher ratio among the towns in this data set?

```
median(Boston$ptratio)
```

```
## [1] 19.05
```

Mediannya 19.05

g. Which census tract of Boston has lowest median value of owner-occupied homes? What are the values of the other predictors for that census tract, and how do those values compare to the overall ranges for those predictors? Comment on your findings.

```
subset(Boston, medv == min(Boston$medv))
```

```
##      crim zn indus chas   nox    rm age    dis rad tax ptratio lstat
medv
## 399 38.3518  0  18.1    0 0.693 5.453 100 1.4896  24 666    20.2 30.59
5
## 406 67.9208  0  18.1    0 0.693 5.683 100 1.4254  24 666    20.2 22.98
5
```

Untuk rumah yang dimiliki pemiliknya, sub-urban ke-399 dan ke-406 memiliki nilai rata-rata terendah.

h. In this data set, how many of the census tract average more than seven rooms per dwelling? More than eight rooms per dwelling? Comment on the census tracts that average more than eight rooms per dwelling.

```
sum(Boston$rm > 7)
```

```
## [1] 64
```

```
sum(Boston$rm > 8)
```

```
## [1] 13
```

```
summary(Boston)
```

```
##      crim      zn      indus      chas
##  Min.   : 0.00632  Min.   : 0.00  Min.   : 0.46  Min.   :0.00000
## 1st Qu.: 0.08205 1st Qu.: 0.00 1st Qu.: 5.19 1st Qu.:0.00000
##  Median : 0.25651  Median : 0.00  Median : 9.69  Median :0.00000
```

```

## Mean : 3.61352 Mean : 11.36 Mean :11.14 Mean :0.06917
## 3rd Qu.: 3.67708 3rd Qu.: 12.50 3rd Qu.:18.10 3rd Qu.:0.00000
## Max. :88.97620 Max. :100.00 Max. :27.74 Max. :1.00000
## nox rm age dis
## Min. :0.3850 Min. :3.561 Min. : 2.90 Min. : 1.130
## 1st Qu.:0.4490 1st Qu.:5.886 1st Qu.: 45.02 1st Qu.: 2.100
## Median :0.5380 Median :6.208 Median : 77.50 Median : 3.207
## Mean :0.5547 Mean :6.285 Mean : 68.57 Mean : 3.795
## 3rd Qu.:0.6240 3rd Qu.:6.623 3rd Qu.: 94.08 3rd Qu.: 5.188
## Max. :0.8710 Max. :8.780 Max. :100.00 Max. :12.127
## rad tax ptratio lstat
## Min. : 1.000 Min. :187.0 Min. :12.60 Min. : 1.73
## 1st Qu.: 4.000 1st Qu.:279.0 1st Qu.:17.40 1st Qu.: 6.95
## Median : 5.000 Median :330.0 Median :19.05 Median :11.36
## Mean : 9.549 Mean :408.2 Mean :18.46 Mean :12.65
## 3rd Qu.:24.000 3rd Qu.:666.0 3rd Qu.:20.20 3rd Qu.:16.95
## Max. :24.000 Max. :711.0 Max. :22.00 Max. :37.97
## medv
## Min. : 5.00
## 1st Qu.:17.02
## Median :21.20
## Mean :22.53
## 3rd Qu.:25.00
## Max. :50.00

```

64 pinggiran kota memiliki rata-rata lebih dari 7 kamar, dan 13 pinggiran kota memiliki rata-rata lebih dari 8 kamar.